



Tech Returners

License: [Attribution-NonCommercial-NoDerivatives 4.0 International](#)



Welcome to the Connecting to a Real Database Lab - TypeScript / PostgreSQL Version

In this lab, you'll be exploring how to connect your Minimalist Book Manager API to work with a PostgreSQL database. You'll also learn about the practice of externalising application configuration/settings. You'll be doing this in the context of an Express API application.




You'll learn and apply the concept of externalising application configuration/settings through interacting with the `.env` file.

You'll also create different `.env` files for various environments and utilise the Dotenv library.



What do I need for this Lab?

You will need:

- To have **forked** and **cloned** the following Starter Minimalist Book Manager API project:
 -  [GitHub - techreturners/lm-code-book-manager-api-ts](https://github.com/techreturners/lm-code-book-manager-api-ts)
 - If you had already forked and cloned this repository as part of an earlier lab or assignment in this series, please feel free to utilise your existing Book Manager API project.
- [PostgreSQL](#) (please see the  **PostgreSQL** section for more information)
- [Postman Desktop Application](#) (please see the  **Postman Desktop Application** section for more information)
- [Google Chrome Web Browser](#)
- [Visual Studio Code](#)

PostgreSQL Primer

This lab assumes that you already have PostgreSQL installed and have working knowledge of starting and stopping a PostgreSQL Server. This lab also assumes you are able to write simple SQL queries.

Postman Desktop Application

This lab assumes that you already have the Postman Desktop Application installed, and that you have some familiarity with HTTP.

1 Explore the Database connection setup

The Minimalist Book Manager API is currently utilising a SQLite In-Memory database. Let's remind ourselves of the clues that indicate this.

🔍 Navigate to the `database.ts` file within the **src/database** directory and look at the code.

```
1 import { Sequelize } from "sequelize";
2
3 // TODO: This should be external config
4 export const sequelize = new Sequelize("sqlite::memory:");
```

🔍 Notice the import from the `sequelize` library. It is possible to manually write all of our database code, including manually writing SQL commands to execute, but it is much **easier** and **safer** to use a library which abstracts over the database. *Sequelize* is one of the most popular JavaScript (and hence TypeScript) libraries for connecting to databases.

💡 There are many similar libraries to choose from, including more modern approaches such as [Prisma](#) which makes powerful use of TypeScript features. However, *Sequelize* is a great choice and, at the time of writing, is used more widely in production thanks to its huge headstart.

🔍 Notice the parameter being passed to Sequelize is listed as `sqlite::memory:`. That parameter is the **database connection string**.

A **connection string** is what it sounds like: a *string* which tells a database library how to connect to a particular database or server. These come in all shapes and sizes, depending on which kind of database we're connecting to. You can find handy reference for many different types of connection string at [ConnectionStrings.com](https://connectionstrings.com)

Specifically this connection string tells us that the Sequelize engine is connecting to a SQLite database running in memory.

🔍 Notice that this connection string is **hard-coded**. If we wanted to change that connection string we would have to update our code and re-build then re-deploy the application 😞

2 Connecting your Minimalist Book Manager API to a PostgreSQL Database - Adding the PostgreSQL Dependency

Since the Minimalist Book Manager API utilises an In-Memory database, the application does not initialise with any books.



An In-Memory database uses temporary memory for data storage; whereas a “real” database persists data using permanent disk storage. This means when the application is shut down, the data held in an In-Memory database is cleared.

If a book is added to the In-Memory database using the POST `api/v1/books` endpoint, that book will only be available for the lifetime that the application is running. It disappears when the app stops running because the in-memory database gets cleared.

To make using the app in development more convenient, we populate the in-memory database with a couple of books on startup.

🔍 If you look in the **src/server.ts** file you'll see we call a function named `populateDummyData()` to populate the database with sample data...

... but only if we're in the dev environment. In production we would not want to mess with the contents of the database every time the app starts! 😱

🔧 In-memory databases are useful tools for development and testing as they reset on each app start. But this approach is clearly inadequate for production. So let's get your Minimalist Book Manager API connected to a real PostgreSQL database.

💡 The steps to configure Sequelize to connect to a database server are similar, regardless of which relational database engine you use. In this lab, we're going to connect to a PostgreSQL database, but connecting to, say, MySQL or MariaDb would be very similar.

🔍 Navigate to your **package.json** file. You should notice that we already include the **sqlite3** dependency:

```
"devDependencies": {  
  ...  
  "sqlite3": "^5.1.4",  
  ...  
}
```

🔍 To use PostgreSQL, we need install a package which Sequelize can use to communicate with a PostgreSQL server.

👉 Open the terminal and ensure you are at the root of your code base (where the `package.json` file is located) and run the following command:

```
npm install -D pg pg-hstore dotenv @types/pg
```

<https://www.npmjs.com/package/pg>

<https://www.npmjs.com/package/dotenv>

This will install the PostgreSQL dependency and the `dotenv` dependency.

Your dependencies should now include `pg`, `pg-hstore` and `dotenv`, as well as some types for `pg`:

```
"devDependencies": {
  "@types/express": "^4.17.15",
  "@types/jest": "^29.2.5",
  "@types/node": "^18.11.18",
  "@types/pg": "^8.6.6",
  "@types/sequelize": "^4.28.14",
  "@types/sqlite3": "^3.1.8",
  "@types/supertest": "^2.0.12",
  "dotenv": "^16.0.3",
  "eslint-config-prettier": "^8.6.0",
  "jest": "^29.3.1",
  "nodemon": "^2.0.20",
  "pg": "^8.8.0",
  "pg-hstore": "^2.3.4",
  "sqlite3": "^5.1.4",
  "supertest": "^6.3.3",
  "ts-jest": "^29.0.3",
  "ts-node": "^10.9.1",
  "typescript": "^4.9.4"
},
```

🔍 Take a look at the **book.ts** file in the **src/models** directory and the **books.ts** file in the **src/services** directory. Notice that neither file specify any SQL. This is the power of the libraries like Sequelize - they *abstract* over the database engine for you so that the library will generate the SQL to query your target database.

🔍 In **book.ts** we define a Sequelize model and then the library takes care of the rest. That model represents a table in your database. This means that we won't have to change our code when we connect to a real database!



📖 This would be a great time to read more about how Models work in the Sequelize documentation: [📄 Model Basics | Sequelize](#)

3 Using Dotenv and setting up the connection settings .env file

It would be great to be able to change our database settings without changing any code. We can do this by moving our settings into an environment file. This allows for secure storage of secrets, such as database user names and passwords, as well as a convenient place to store application settings outside of code.

👉 Firstly create a file called **.env.dev** at the root of your project. (The same location as the package.json file).

NOTE: The dot at the start of that filename is **important**.

👉 It's **.env.dev** not **env.dev**

🔍 Add the following contents to your **.env.dev** file:

```
NODE_ENV=dev
PORT=3000
DB_NAME=sqlite::memory:
DB_USERNAME=
DB_PASSWORD=
DB_HOST=localhost
DB_PORT=
DB_DIALECT=sqlite
```

This configuration file is going to hold our database connection settings.

🔍 Next we need to update **server.ts** to use dotenv.

👉 Place the following line at the very top of your **server.ts**

```
import * as dotenv from "dotenv";
const environment = process.env.NODE_ENV || "dev";
dotenv.config({ path: `./.env.${environment}` });
```

So your **server.ts** file should now look similar to this:

```
import * as dotenv from "dotenv";
const environment = process.env.NODE_ENV || "dev";
dotenv.config({ path: `./.env.${environment}` });

// ! NOTE ! The dotenv.config needs to happen BEFORE we do
these imports, as they may rely on the config...

import { app } from "./app";
import { populateDummyData } from "./database/database_seed";

const PORT = 3000;

console.log(`🌍 Running in ${environment} environment`);

app.listen(PORT, () => {
  console.log(`🚂 Express started on port ${PORT}`);

  // Seed the database with some data
  if (environment === "dev") {
    populateDummyData();
  }
});
```

🔍 Let's try using some of the properties we put in our `.env` file. On line 9 of the code where it lists 3000 as the port, try replacing the 3000 number with `process.env.PORT`.

```
import * as dotenv from "dotenv";
const environment = process.env.NODE_ENV || "dev";
dotenv.config({ path: `./.env.${environment}` });

import { app } from "./app";
import { populateDummyData } from "./database/database_seed";

const PORT = process.env.PORT;

console.log(`🌍 Running in ${environment} environment`);

app.listen(PORT, () => {
  console.log(`🚀 Express started on port ${PORT}`);

  // Seed the database with some data
  if (environment === "dev") {
    populateDummyData();
  }
});
```

🎉 You've just used your first piece of external configuration!

Next, let's use this to connect to a real database.

🔍 Open up the **database.ts** file located in the **src/database** directory. It is currently looks like this:

```
import { Sequelize } from "sequelize";

// TODO: This should be external config
export const sequelize = new Sequelize("sqlite::memory:");
```

🔍 Update this file to use the properties you've set up in your **.env.dev** file:

```
import { Dialect, Sequelize } from "sequelize";

export let sequelize = new Sequelize("sqlite::memory:");

if (process.env.NODE_ENV !== 'test') {
  sequelize = new Sequelize(process.env.DB_NAME ??
    'MISSING_DB_NAME_CONFIG',
    process.env.DB_USERNAME ??
    'MISSING_DB_USERNAME_CONFIG',
    process.env.DB_PASSWORD ??
    'MISSING_DB_PASSWORD_CONFIG', {
    host: process.env.DB_HOST ?? 'MISSING_DB_HOST_CONFIG',
    port: parseInt(process.env.DB_PORT as string) ??
    'MISSING_DB_PORT_CONFIG',
    dialect: (process.env.DB_DIALECT as Dialect) ??
    'postgres',
  });
}
```

🤔 Let's think a bit more about some of these changes...

🔍 By default, we use an In-Memory database. However, if we're NOT in a test environment, then we replace that connection with whatever is configured for our environment.

🔍 Notice that TypeScript forces us to provide defaults, as there's no guarantee that `process.env.DB_NAME` will exist. In this case, it's sensible to provide some obvious default that will indicate exactly what the problem is if we see it pop up in an error log.

🔍 `Dialect` is a type that is part of Sequelize. We have to indicate to TypeScript that whatever we put in here ought to be one of the defined dialects that Sequelize speaks. We'll default to `postgres` for now as we're planning to connect to PostgreSQL anyway.

👉 See here for more on `Dialect` :

<https://sequelize.org/docs/v6/other-topics/dialect-specific-things/#postgres>

! Notice in the documentation above that the `postgres` dialect depends on a library called `"pg"` to speak to PostgreSQL.

👉 Install `"pg"` so Sequelize can communicate with PostgreSQL:

```
npm install pg
```

4 Connecting your Minimalist Book Manager API to a PostgreSQL Database - Adding the Database Connection String to the .env file

To find out the database connection string to your PostgreSQL database, you will need to check the port that the database is using on your local machine.

🔍 Start PostgreSQL Server.

🔍 Open up your terminal / command prompt and run the following to utilise the PostgreSQL client.

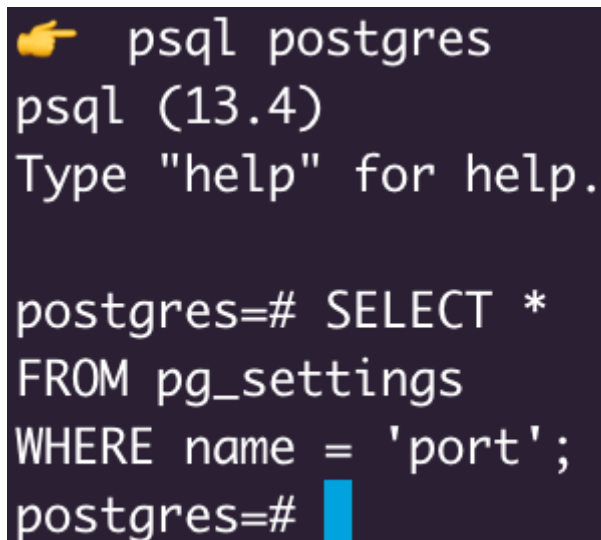
```
psql postgres
```

💡 You may choose to utilise [pgAdmin](#) instead of the terminal if you prefer.

👉 At the postgres=# command prompt, type \l to list out all the databases.

🔍 Run the following statement to check which port PostgreSQL is running on. The port will usually be 5432 , but you may have PostgreSQL running a different port.

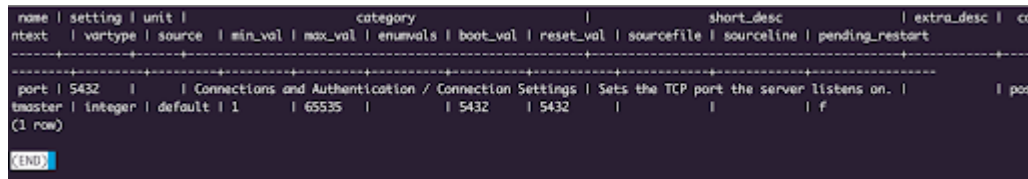
```
SELECT * FROM pg_settings WHERE name = 'port';
```



```
👉 psql postgres
psql (13.4)
Type "help" for help.

postgres=# SELECT *
FROM pg_settings
WHERE name = 'port';
postgres=#
```

Running the SQL statement to find out which port PostgreSQL is running on



```
name | setting | unit | category | short_desc | extra_desc | co
ntext | vartype | source | min_val | max_val | enumvals | boot_val | reset_val | sourcefile | sourceline | pending_restart
-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----
port | 5432 | | Connections and Authentication / Connection Settings | Sets the TCP port the server listens on. | | pos
tmaster | integer | default | 1 | 65535 | | 5432 | 5432 | | | f
(1 row)
```

<END>

This is an example output from the SQL SELECT query which indicates the port that PostgreSQL is running on is 5432. Your port may differ.

👉 Let's create a new database called bookshop for our app to use. You may already have this database set up from a previous lab or assignment. If not, run the following SQL statement:

```
CREATE DATABASE bookshop ;
```

👉 Type \1 to list out all of the databases to double-check that the bookshop database has been created successfully. Notice the username listed in the Owner

column against the bookshop database.

```
12:04:55
👉 psql postgres
psql (13.4)
Type "help" for help.

postgres=# \l

               List of databases
  Name      | Owner   | Encoding | Collate | Ctype | Access privileges
-----+-----+-----+-----+-----+-----
databaselab | kimdiep | UTF8     | C       | C     |
postgres    | kimdiep | UTF8     | C       | C     |
template0   | kimdiep | UTF8     | C       | C     | =c/kimdiep +
            |         |          |         |       | kimdiep=CTc/kimdiep
template1   | kimdiep | UTF8     | C       | C     | =c/kimdiep +
            |         |          |         |       | kimdiep=CTc/kimdiep
(4 rows)

postgres=# CREATE DATABASE bookshop;
CREATE DATABASE
postgres=# \l

               List of databases
  Name      | Owner   | Encoding | Collate | Ctype | Access privileges
-----+-----+-----+-----+-----+-----
bookshop    | kimdiep | UTF8     | C       | C     |
databaselab | kimdiep | UTF8     | C       | C     |
postgres    | kimdiep | UTF8     | C       | C     |
template0   | kimdiep | UTF8     | C       | C     | =c/kimdiep +
            |         |          |         |       | kimdiep=CTc/kimdiep
template1   | kimdiep | UTF8     | C       | C     | =c/kimdiep +
            |         |          |         |       | kimdiep=CTc/kimdiep
(5 rows)
```

CREATE DATABASE bookshop; and \l to check that the database has been created successfully

👉 It would be good to create a specific user for our app to use. Create a user: (If you like you can replace the password with another.)

```
CREATE USER bookshopuser WITH PASSWORD
'super-secret-password' ;
```

👉 And let's give that user control over the bookshop database:

```
GRANT ALL PRIVILEGES ON DATABASE bookshop TO bookshopuser ;
```




Normally it is good practice to be more restrictive with permissions than this. For example, our API doesn't *really* need to create/delete entire **tables**, just data in those tables, so we shouldn't really give it ALL privileges. However, for the sake of ease, we're being lazy in this instance 🙄 Just be aware of this in the real world!

👉 Now use the \c command to connect to the bookshop database.

```
\c bookshop;
```

👉 And then create a table in the bookshop called Books

```
CREATE TABLE IF NOT EXISTS "Books" ("bookId" INTEGER PRIMARY KEY, "title" VARCHAR(255) NOT NULL, "author" VARCHAR(255) NOT NULL, "description" VARCHAR(255) NOT NULL);
```

👉 Add access to the tables:

```
GRANT ALL ON TABLE public."Books" TO bookshopuser;
```

👉 We can now insert a book into the Books table

```
INSERT INTO "Books" VALUES (1, 'Fantastic Mr. Fox', 'Roald Dahl', 'This book concerns a fox who is fantastic.');
```

🔍 Now let's change the database settings in the **.env** file to enable your Book Manager API to connect to your PostgreSQL database instead of the SQLite In-Memory database.

🔍 Add a new file called `.env.production` next to the `.env.dev` file. In here, place your production configuration:

```
NODE_ENV=production
PORT=3000
DB_NAME=bookshop
DB_USERNAME=root // ! replace with your db username
DB_PASSWORD=my-secret-pw // ! replace with your db password
DB_HOST=127.0.0.1
DB_PORT=5432
DB_DIALECT=postgres
```

N.B. If your port was different to 5432 please change the port number accordingly, substituting 5432 for the port number where PostgreSQL is running on your machine. Also you'll need to make sure the **password is correct for your database user**.

It's also common to define a default `.env` rather than explicitly naming them after their environments.

👉 On Mac/Linux, in `package.json` you can add this new script to start the app in production mode:

```
"scripts": {  
  "start": "nodemon src/server.ts",  
  "start-prod": "NODE_ENV=production nodemon src/server.ts",  
  "test": "jest --watch"  
},
```

👉 On Windows, it's a little trickier to temporarily alter environment variables like this. On some versions of Windows, you may find that this works in place of the above command:

```
SET NODE_ENV=production & nodemon src/server.ts
```

On some versions of Windows, you may find that this works: 🔍 Try stopping your API and restarting by using the command `npm run start-prod`

🔍 Try hitting your API to GET all books - you should see books from your PostgreSQL database now instead of the in-memory SQLite database.

HELP!

If you get connection errors, you may find it easier to replace your database connection code with a URL-based connection format:

```
const connString =  
"postgres://bookshopuser:super-secret-password@127.0.0.1:5432/  
bookshop";
```

**// TODO: Replace the above string with a string built from
YOUR `process.env.DB_USERNAME` etc variables**

```
sequelize = new Sequelize(connString);
```

🎉 Well done! You just successfully swapped out the connection to the SQLite In-Memory Database with the connection string to the PostgreSQL Database!

The next step is to see this in action.

5 Use the Postman Desktop Application to interact with the Minimalist Book Manager API

Since the Minimalist Book Manager API is now connected to a real database, the data will persist in database storage even when the Book Manager API shuts down.

🔍 Try adding some data into your Book table. Using the Postman Desktop Application, you can make a POST request to the <http://localhost:3000/api/v1/books> endpoint to add a book.

Here are some example JavaScript Object Notation (JSON) formats for some books you can post inside the HTTP request body.

```
{
  "bookId": 2,
  "title": "You Don't Know JS Yet: Get Started",
  "author": "Kyle Simpson",
  "description": "Learn JS with Kyle Simpson"
```

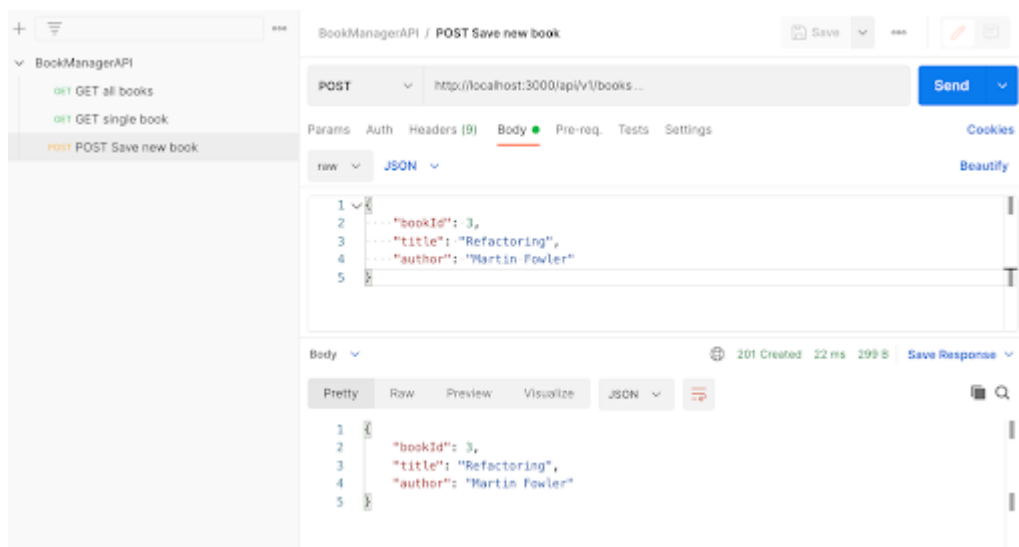
```

}

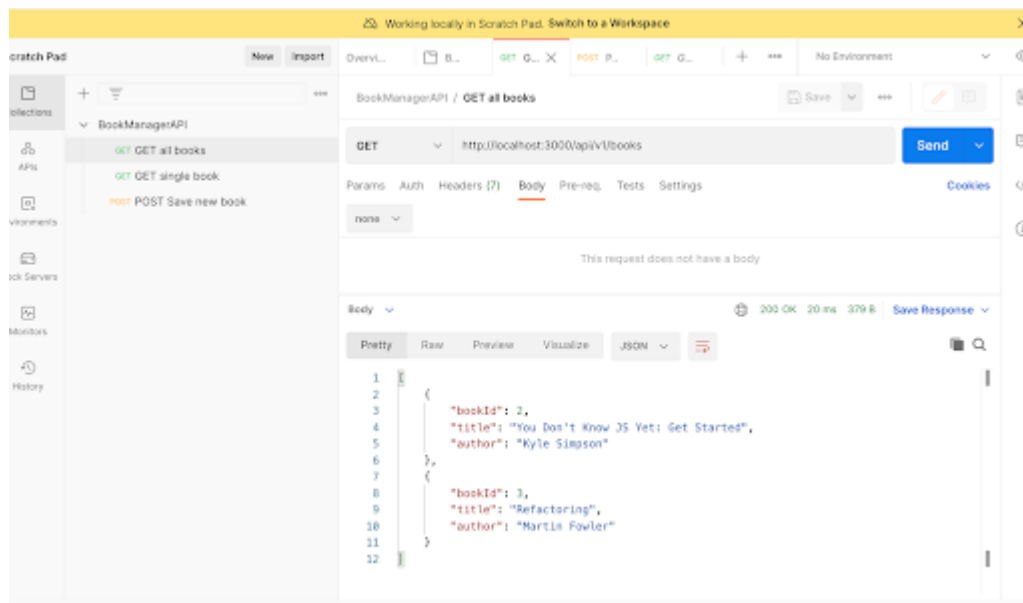
{
  "bookId": 3,
  "title": "Refactoring",
  "author": "Martin Fowler",
  "description": "How to refactor"
}


{
  "bookId": 4,
  "title": "The Shop Before Life",
  "author": "Neil Hughes",
  "description": "Before being born, each person must visit
the magical Shop Before Life, where they choose what kind of
person they will become down on Earth..."
}

```




🔍 Use the Postman Desktop Application to make a GET request to the <http://localhost:3000/api/v1/books> endpoint to get all books from your PostgreSQL database. You should be able to see something that looks similar to this. The books are being retrieved from the PostgreSQL database.



 **Experiment!** Using the PostgreSQL Client - either in the terminal or through pgAdmin - add a new book into the Books table using an INSERT SQL statement. What do you think the GET request to the <http://localhost:3000/api/v1/books> endpoint will return?

 **Top Tip:** You need to conduct a new GET request to the endpoint.

 **Experiment!** Stop your Minimalist Book Manager API application. Using the PostgreSQL Client (PostgreSQL Monitor) and ensuring you are using the bookshop database, run the following SQL statement. What do you expect to see? Will the books data persist or will the data disappear?

```
SELECT * FROM "Books";
```

6 .env files and git

🔍 Notice that the dev version of the env file and the prod version of the env file are **not being included within version control**. This is managed by the **.gitignore**.

! This is very important! For security purposes we don't want to include configuration files in version control, as otherwise we make our secrets *permanently part of the codebase* 😱😱

License: [Attribution-NonCommercial-NoDerivatives 4.0 International](#)